

Music Production Assisted by Machine Learning

Nianze Liu
New York University

This article overviews recent progress in machine learning as a creative assistance in music production. Specifically, two innovative machine learning models, NSynth and Coconet, are introduced and discussed in the context of music composition, music production, and music live performance.

Keywords: Computer Music, Machine Learning

Introduction

Since the invention of computer, music industry has hugely been influenced. Nowadays, we can find tons of music pieces in the digital format: some might be played by acoustic instruments and recorded down, others may be generated by synthesizers, but no matter how the sound is generated, the music waves end up into a sequence of digital representations, which turns out to be a good data source for machine to learn from. With the hardware becomes more and more powerful in the recent years, more complicated machine learning models become feasible for providing helpful assistance in various areas, such as automatic driving and speech recognition. Music production is no exception under this trend. By the help of new machine learning models, we might train the computer to learn the musical master's composition techniques to help generate harmonies given a piece of melody we write; it is also possible to synthesize new sound that might be very difficult to create by the traditional synthesizer-based methods.

In this article, two machine learning models are discussed. We will go over their basic structures to see how they work. After understanding the abilities of these innovative models, we will discuss what might be changed by introducing these techniques into the current music composition, music production and even live music performance, and what might be achieved in the future.

Computer Assisted Composition

How Computer Generates Music

Traditionally, when computer generates music, the most natural way is to do the job in a sequence, under an orderly manner, because when computer generates music, the only thing it understands is just numbers, or in our context, statistics and probability - it does not feel like humans do. Given a note, such as C4, and another note, say E4, the computer has to calculate all the possibilities of the next note, and return the most possible notes to fill in the next position. If

the model sees 90 percent of time that C4, E4 and G4 always show together during training, the likelihood of generating a G4 given the context of C4 and E4 is then higher than other options – this higher possibility is stored in the model in the form of network weight and we can say that the model 'learns' the C major triad now.

Under this sequential generation process, traditionally, before determining the very first note of music, the computer has to have a complete view of the whole piece of music, since as long as the previous note is determined, there's no way back to refine it anymore – the note has to be generated in a sequence in one iteration. This orderly generation requirement is very strict, which leads to the creation of a new machine learning model named 'Coconet'.

Coconet

The machine learning model Coconet is basically a convolutional neural network but is integrated with some creative process. The details of the model are described in Huang, Cozijmans, Roberts, Courville, and Eck (2019). Here we will just go over the basics of the model structure and see how it works to generate four voices counter points and harmonize any given input melody.

The biggest idea behind the model is to mimic the process of human composition. The key observation here is that, when human beings compose, the composer seldom writes the music in sequence. Instead, human composer tends to rewrite the same piece back and forth over and over again in order to get the better refined result that finally meets his or her standard. In other words, the composition is typically done in an orderless manner. To achieve orderless composition by computer, two techniques are created: Orderless modeling and Gibbs sampling.

The basic idea is to use Gibbs sampling to generate music in multiple iterations, and in every iteration, part of music scores is erased, so that we can use Orderless modeling to come up with the independent conditional probabilities for missing notes. For the same example of C4 E4 G4 triad

above, if we mask the E4 in the middle, the model should be able to generate E4 given the context of C4 in the first position and G4 in the 3rd position. In math, the model is basically computing:

$$P(x_2|x_1, x_3), \text{ where } x_n \text{ means score at beat } n$$

Since the model doesn't know which note will be erased, in orderless modeling, all possible generation sequences have to be modeled. In the example of 3 notes, there are $3! = 6$ possible permutations.

Gibbs sampling, on the other hand, is just a fancy word for phrasing our repeated iterative music generation/modification in multiple passes, and orderless modeling guarantees that no matter which note is randomly erased by a mask in a random iteration, the model is able to generate a conditional distribution for the missing score. Figure 1 below shows this multi-iteration music notes generation and erasing during the training phase.




Figure 1. Music generation in multiple iterations

The result of Coconet is very interesting: given a partial piece of scores for a fixed length, the model is able to complete the music based on how it gets trained. On 2019 March 20, in order to celebrate J.S. Bach's 334th birthday, Google released a Doodle powered by this coconet model, which learns from "a dataset of 306 chorale harmonizations by Bach", according to Magenta Blog post.

Computer Assisted Synthesis

The output of coconet is abstract music notes, which can not be heard directly and functions similarly to MIDI in that we can use synthesizers to generate different timbre for the generated note. Traditional hardware and software synthesizers generate audio by pre-designed oscillators and wavetables. By taking advantage of machine learning, however, there comes a new way to generate audio: by using deep neural networks to extract multi-dimensional sound representations from training data set, and then applying high level modification over these representations to generate original new sound that could be very hard to generate in the tradi-

tional way. The NSynth model introduced in Engel et al. (2017) follows this thought and achieves a remarkable result.

NSynth

The basic structure of NSynth is a fully convolutional neural network, which is pretty similar to the structure of Wavenet introduced in Oord et al. (2016). The goal of NSynth is to model musical qualities other than pitch, so that combined with higher-level note information, audio could be generated and modified easily. In math, we can express this process as:

$$P(\text{audio}) = P(\text{audio}|\text{note})P(\text{note})$$

As shown in Figure 2, the process of training the NSynth model involves an autoencoder and a decoder. The encoder, consisting of 30 layers of computation, analyzes the input audio chunk and extracts out temporal embeddings, which is a vector of 16 dimensions sampled from every 512 samples. This vector represents the musical qualities we want to model in a very abstract way. The decoder, on the other hand, is the same as WaveNet decoder. After applying a 1×1 convolution between the input audio chunk as well as the embeddings generated from the encoder, we feed this result as a bias to each of decoder's 30 layers.




Figure 2. Training workflow for NSynth

The key in NSynth is the 16-dimension temporal embeddings vector. With this vector, we can do all kinds of interesting operations to create novel sounds, such as timestretching, interpolation, and composition. Since the operation is applied to this high level vectors instead of the audio waves itself, the effect of operation sounds differently from the traditional sound effects. For example, by adding two temporal embeddings together, say a vector representing the sound of a flute and another one representing the sound of a dog barking, the result is a single merged sound that lies between the two source sounds, instead of two sounds played at the same.

Future Possibilities

Combining the two models mentioned above, we can already see some future possibilities of music composition and music production. By the help of Coconet, composers might get more inspiration of possible harmonizations, so that they may compose new music in a more efficient way. With

the musical notes at hand, no matter they are generated by the machine learning models or hand-crafted by human beings, musicians are able to applying noval sound effects to the notes thanks to models such as NSynth. Moreover, the NSynth decoder is able to generate sound on the fly, so it is possible to apply this synthesis techniques into live performance. With all these machine learning models assisting the creative process in more phases of music production, musicians are able to explore a wider areas that could never be imagined before.

References

- Engel, J., Resnick, C., Roberts, A., Dieleman, S., Norouzi, M., Eck, D., & Simonyan, K. (2017). Neural audio synthesis of musical notes with wavenet autoencoders. In *Proceedings of the 34th international conference on machine learning-volume 70* (pp. 1068–1077). JMLR. org.
- Huang, C.-Z. A., Cooijmans, T., Roberts, A., Courville, A., & Eck, D. (2019). Counterpoint by convolution. *arXiv preprint arXiv:1903.07227*.
- Oord, A. v. d., Dieleman, S., Zen, H., Simonyan, K., Vinyals, O., Graves, A., ... Kavukcuoglu, K. (2016). Wavenet: A generative model for raw audio. *arXiv preprint arXiv:1609.03499*.